

# Visual Recognition using Deep Learning

---

## Lab1 report

---

**110550130 劉秉驊**

### My github

1. [https://github.com/Potato-TW/visual\\_dl/tree/main](https://github.com/Potato-TW/visual_dl/tree/main)

### Introduction

In this homework, we use `resnext101_32x8d.fb_swsl_ig1b_ft_in1k` with pre-trained weights from timm to classify 100 kinds of plants.

We implement several methods to raise accuracy and suppress loss on valid like adjustment of hyperparameters, different transformation of images, and even freezing layers in model.

We hope our fine-tuned model can reach higher performance than official training.

### Method

1. Choosing model

Our model is `resnext101_32x8d.fb_swsl_ig1b_ft_in1k`.

In timm leaderboard of image classification, this model gets 90% in top-5 accuracy with only 88M parameters.

It's really powerful.

model	img_size	avg_top1	avg_top5	imagenet_top1	imagenet_top5	real_top1	real_top5	v2_top1	v2_top5
resnetrs200.tf_in1k	256	66.82	88.95	83.89	96.76	88.29	97.94	73.39	91.28
resnetrs200.tf_in1k	320	65.66	88.46	84.45	97.87	88.75	98.10	74.19	91.86
resnext101_32x8d.fb_sswl_ig1b_ft_in1k	224	76.11	98.08	94.28	97.17	88.76	98.14	75.46	92.68
resnext101_32x8d.fb_wsl_ig1b_ft_in1k	224	75.09	98.09	82.72	96.62	88.17	97.86	73.71	92.18
resnext101_32x8d.tv_in1k	224	64.62	79.36	82.83	96.23	87.32	97.56	72.24	90.16
resnext101_32x8d.fb_ssl_yfcc100m_ft_in1k	224	64.61	88.98	81.62	96.84	86.88	97.47	71.43	90.31
resnext101_32x8d.tv2_in1k	176	64.22	79.28	81.99	95.72	86.65	97.28	71.29	89.31
resnext101_32x8d.tv_in1k	224	60.87	77.83	79.31	94.52	85.22	96.44	67.76	87.35
seresnext101_64x4d.gluon_in1k	224	63.24	78.63	88.89	95.31	85.95	96.98	70.47	89.18
resnetrs152.tf_in1k	256	64.98	88.84	82.89	96.29	87.69	97.52	71.84	90.28
resnetrs152.tf_in1k	320	64.85	79.75	83.71	96.61	88.26	97.73	72.98	91.21
resnext101_64x4d.tv_in1k	224	65.17	88.18	82.98	96.25	87.72	97.58	72.28	90.61

Model Details

- Model Type: Image classification / feature backbone
- Model Stats:
  - Params (M): 88.8
  - GMACs: 16.5
  - Activations (M): 31.2
  - Image size: 224 x 224

2. Image transformation  
Here is model configure. Model takes 224 \* 224 images as input.

		0
url	<a href="https://download.pytorch.org/models/resnext101_32x8d-8ba56ff5.pth">https://download.pytorch.org/models/resnext101_32x8d-8ba56ff5.pth</a>	
hf_hub_id	timm/	
custom_load	False	
input_size	(3, 224, 224)	
fixed_input_size	False	
interpolation	bilinear	
crop_pct	0.875	
crop_mode	center	
mean	(0.485, 0.456, 0.406)	
std	(0.229, 0.224, 0.225)	
num_classes	1000	
pool_size	(7, 7)	
first_conv	conv1	
classifier	fc	
license	bsd-3-clause	
origin_url	<a href="https://github.com/pytorch/vision">https://github.com/pytorch/vision</a>	

Then we implement several methods on training data, but only crop validation data in center to  $224 * 224$ :

1. Random crop  $224 * 224$
2. Horizontal, vertical flip
3. Rotation 30 degree

## 4. Gaussian blurring

```
visual_dl [WSL: Ubuntu] - 9.ipynb

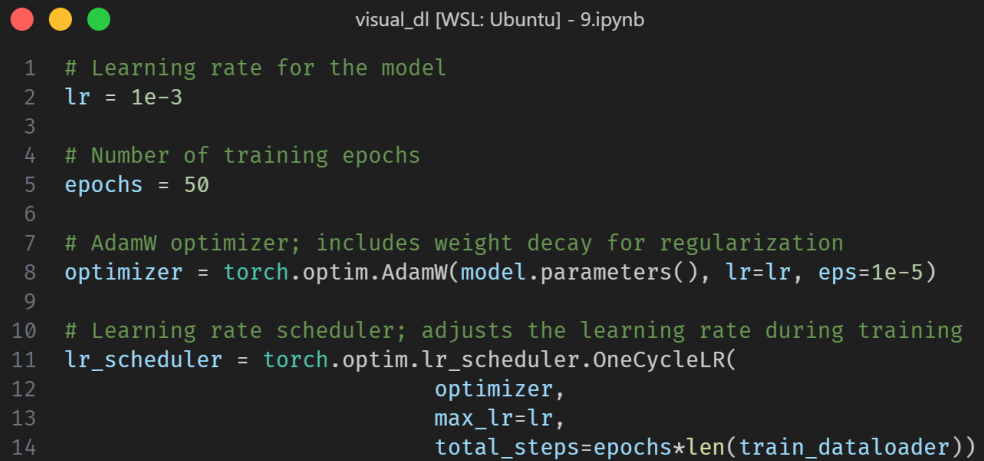
1  import torchvision.transforms.v2 as T
2
3  train_tfms = T.Compose([
4      T.RandomResizedCrop(224),
5      T.RandomHorizontalFlip(p=0.8),
6      T.RandomVerticalFlip(p=0.3),
7
8      T.RandomRotation(30),
9      T.RandomApply([T.GaussianBlur(3)], p=0.5),
10     T.ToImage(),
11     T.ToDtype(torch.float32, scale=True),
12     T.Normalize(*norm_stats)
13 ])
14
15
16 valid_tfms = T.Compose([
17     T.Resize(train_sz, antialias=True),
18     T.CenterCrop(int(train_sz * 0.875)),
19     T.ToImage(),
20     T.ToDtype(torch.float32, scale=True),
21     T.Normalize(*norm_stats)
22 ])
```



## 3. Hyperparameter

1. batch size: 96  
(training on 24G vram of L4 GPU)
2. learning rate: 0.003

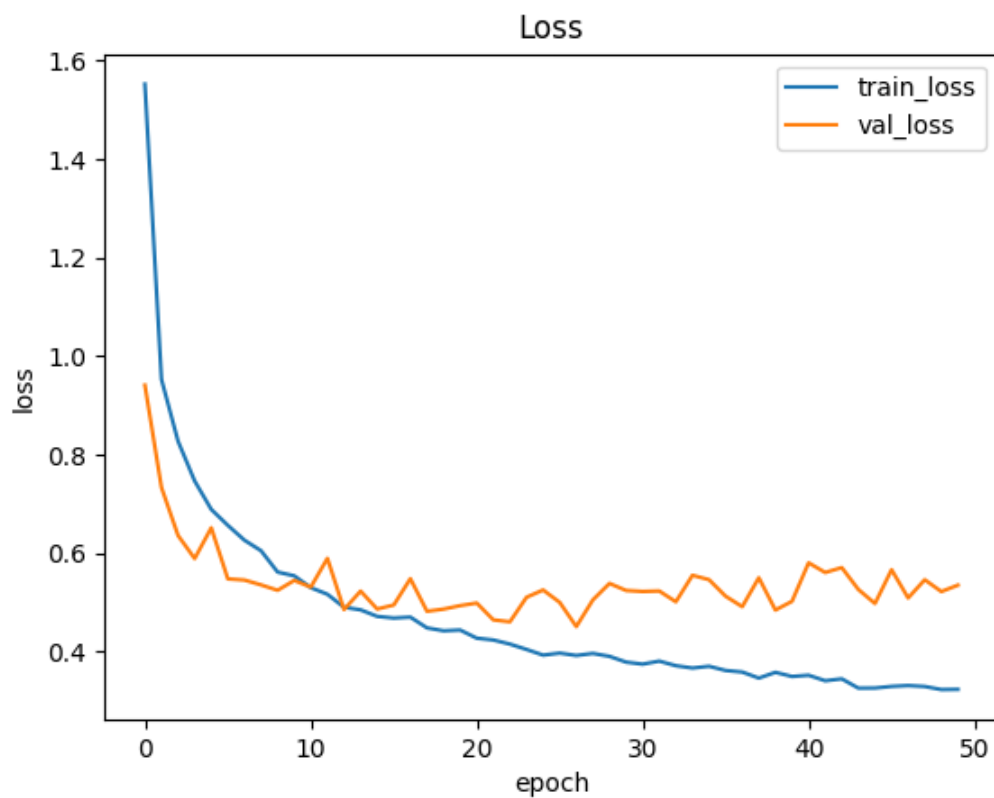
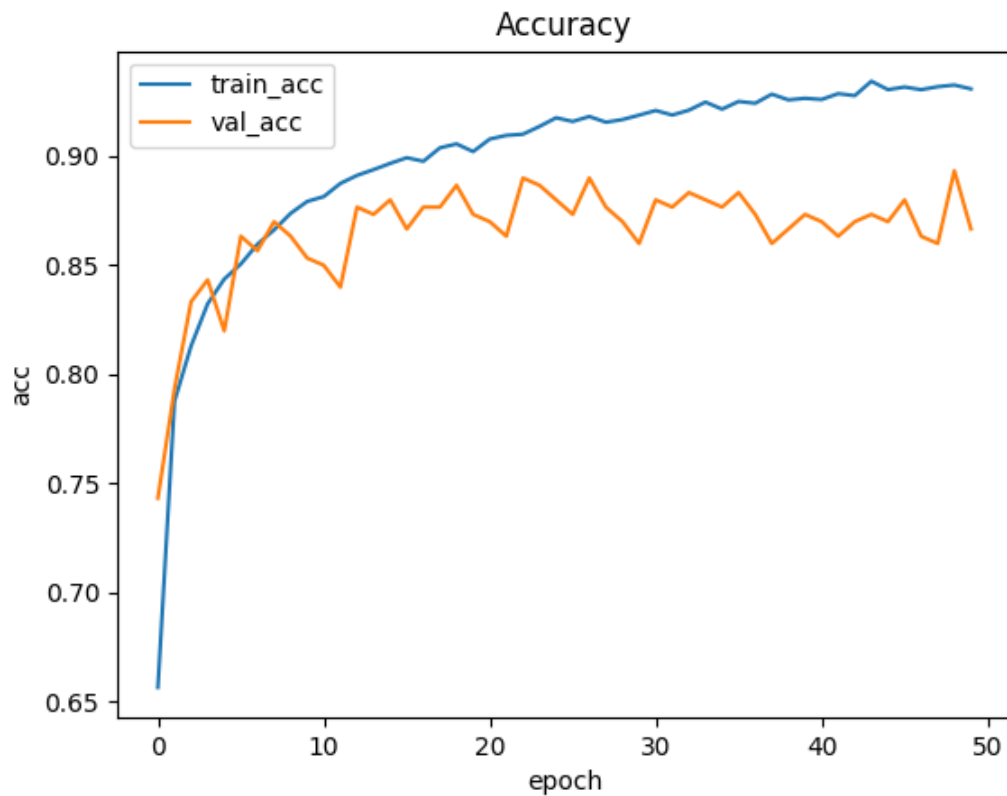
3. epochs: 50

A screenshot of a Jupyter Notebook cell with a dark background. The cell title is 'visual\_dl [WSL: Ubuntu] - 9.ipynb'. The code is written in green and blue text, showing the configuration of a PyTorch model's training process. It includes comments for learning rate, number of epochs, optimizer (AdamW), and a learning rate scheduler (OneCycleLR).

```
1 # Learning rate for the model
2 lr = 1e-3
3
4 # Number of training epochs
5 epochs = 50
6
7 # AdamW optimizer; includes weight decay for regularization
8 optimizer = torch.optim.AdamW(model.parameters(), lr=lr, eps=1e-5)
9
10 # Learning rate scheduler; adjusts the learning rate during training
11 lr_scheduler = torch.optim.lr_scheduler.OneCycleLR(
12     optimizer,
13     max_lr=lr,
14     total_steps=epochs*len(train_dataloader))
```

## Results

1. After training, we can see the trend of accuracy and loss below.  
When in 48<sup>th</sup> epoch, we get the lowest loss 0.155 on training and 0.543 on validation, and meanwhile we get 0.88 of accuracy on validation.  
Finally, we obtain 0.94 of accuracy on test case.



2. At the same time, we also capture the 49<sup>th</sup> epoch of highest accuracy.

0.153 on training loss

0.552 on validation loss

0.887 on validation accuracy

As the result, on test case, we surprisingly get as the same score as we choose lowest loss.

We think the performance of lowest loss and highest accuracy is the same, which even they have different accuracy on validation data.

## Additional experiments

### 1. Freeze sensitive layers

First we need to define sensitive layers.

We use validation data to compute the loss of each layer in the original model.

```
[('layer4.1.conv1.weight', 22.528321981430054),
 ('layer4.2.conv3.weight', 13.603535413742065),
 ('layer4.0.downsample.0.weight', 9.315202116966248),
 ('layer4.2.conv1.weight', 7.713109374046326),
 ('fc.weight', 7.454432964324951),
 ('layer4.0.conv3.weight', 6.771357178688049),
 ('layer4.0.conv1.weight', 6.156463623046875),
 ('layer4.1.conv2.weight', 5.513914704322815),
 ('layer4.0.conv2.weight', 5.070958375930786),
 ('layer3.8.conv1.weight', 5.070551514625549),
 ('layer3.0.conv1.weight', 5.056488037109375),
 ('layer2.0.conv1.weight', 4.999824523925781),
 ('layer4.0.downsample.1.bias', 4.8650208711624146),
 ('layer2.1.conv2.weight', 4.862960934638977),
 ('layer4.2.bn1.bias', 4.816706299781799),
 ('layer3.10.conv2.weight', 4.806696653366089),
 ('layer3.19.conv1.weight', 4.786763548851013),
 ('layer4.0.bn1.weight', 4.784423828125),
 ('layer4.0.bn3.bias', 4.7754669189453125),
 ('layer2.0.conv3.weight', 4.714258670806885),
 ('layer3.5.conv1.weight', 4.7090911865234375),
 ('layer3.1.conv3.weight', 4.708287477493286),
 ('layer3.12.conv1.weight', 4.692952394485474),
 ('layer3.0.bn1.bias', 4.682795286178589),
 ('layer3.0.bn1.weight', 4.67974853515625),
 ('layer3.4.conv1.weight', 4.677983641624451),
 ('layer3.20.conv1.weight', 4.677922487258911),
 ('layer3.16.bn3.bias', 4.677424073219299),
 ('layer3.1.bn3.bias', 4.677154541015625),
 ('layer3.15.conv1.weight', 4.675959348678589),
```

We

freeze the top 2 sensitive layers, expecting this model can keep the capability of learning images.

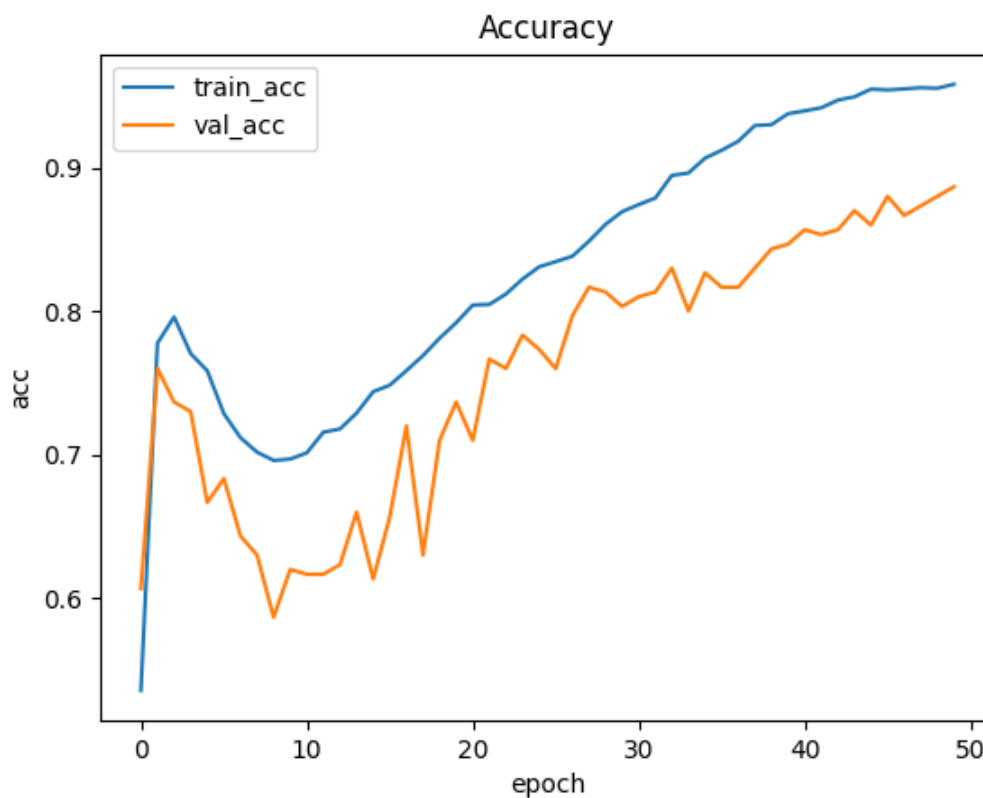
Don't be easily affected by training data.

```
visual_dl [WSL: Ubuntu] - 8.ipynb

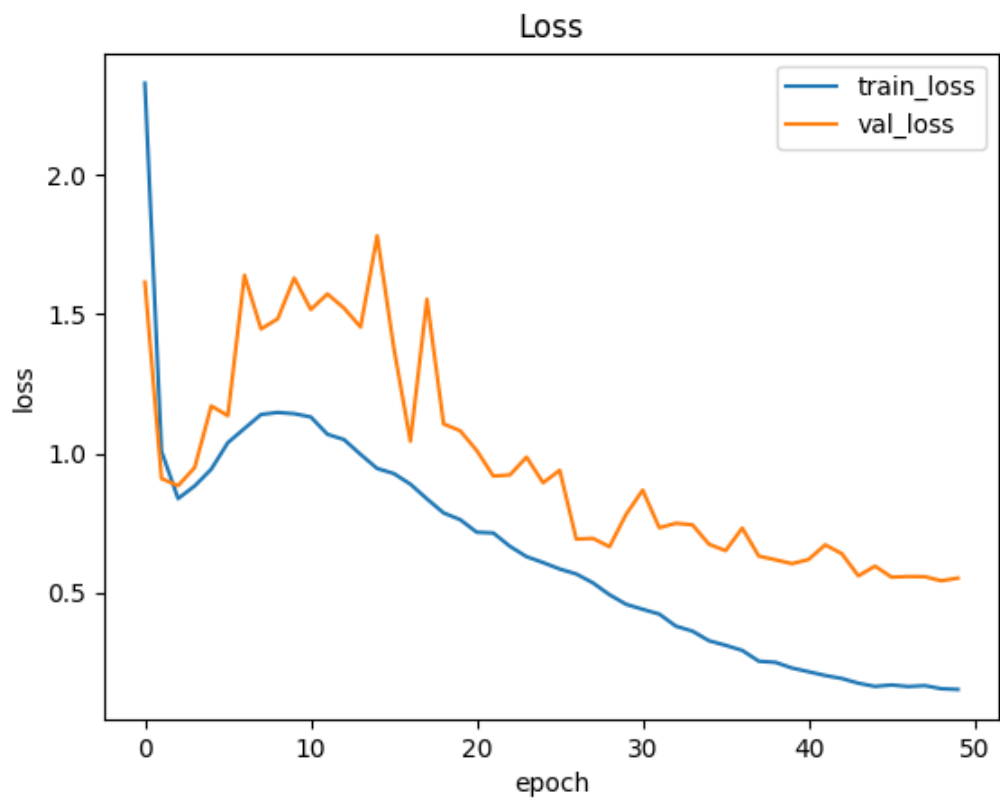
1 protected_layers = [model.layer4[1].conv1, model.layer4[2].conv3]
2
3 # 凍結指定層的參數更新
4 for layer in protected_layers:
5     for param in layer.parameters():
6         param.requires_grad_(False) # 關閉梯度計算
7         param.detach_() # 斷開計算圖(可選強化操作)
8
9 # 驗證凍結狀態
10 frozen_params = [
11     (name, param.requires_grad)
12     for name, param in model.named_parameters()
13     if any(layer in name for layer in ['layer4.1.conv1', 'layer4.2.conv3'])
14 ]
15
16 print(pd.DataFrame(frozen_params, columns=["Parameter", "Trainable"]))
```

We can see the result below.

Accuracy and loss move up then down sharply. But this model can still reach nice accuracy compared with non freezed model, and it even tries to reach lower loss.







We can say this freeze method gets better results than we think.

- 2. Set large learning rate on sensitive layers  
Opposite to freezing method, instead of just freezing layers, we set higher learning rate on those layers, hoping they can learn faster than insensitive layers.

sensitive layer	Top 1	Top 2
learning rate	0.003	0.0008
weight decay	0.005	0.007

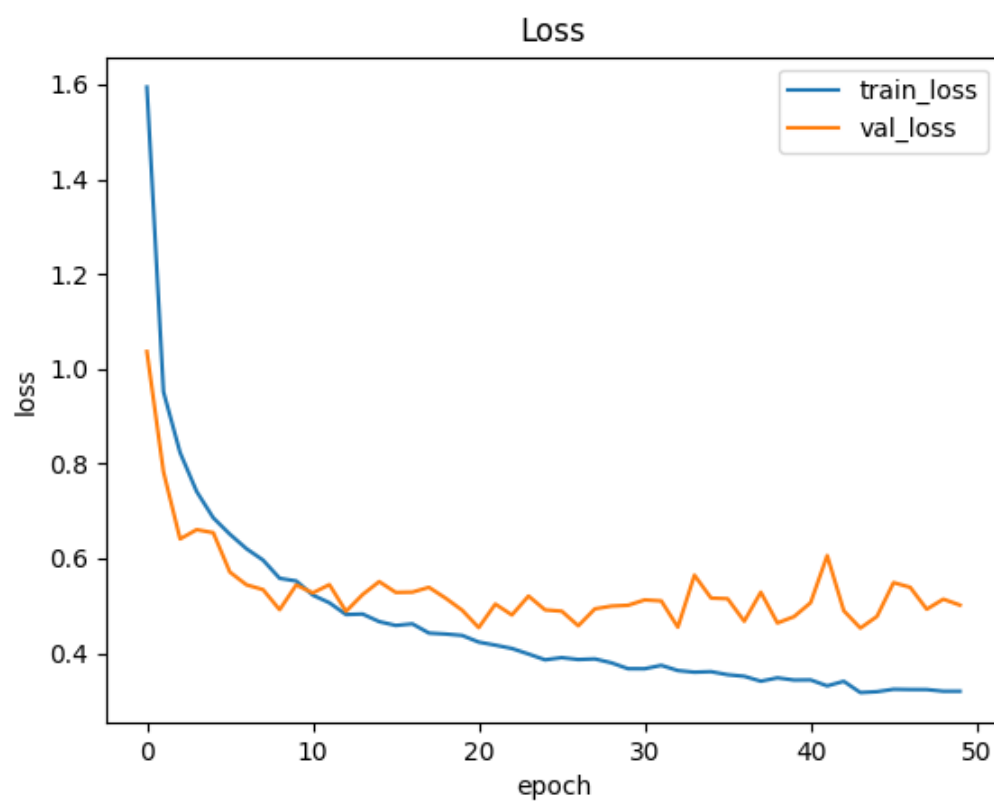
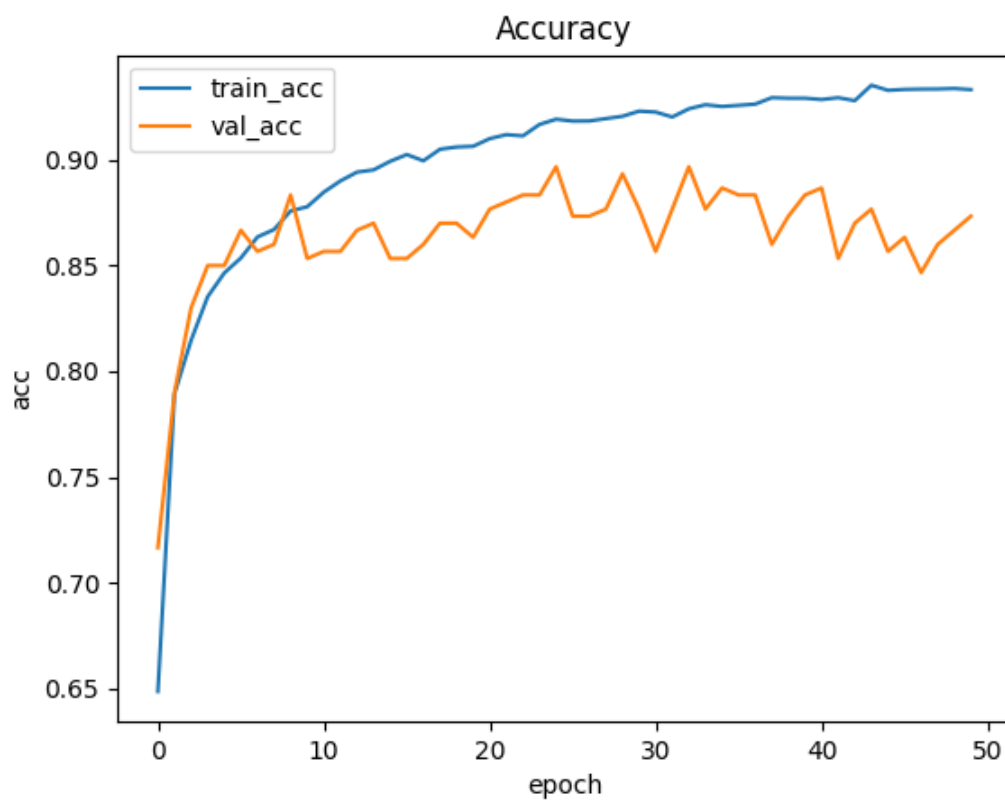


```
visual_dl [WSL: Ubuntu] - 7.py

1  # 重構優化器配置
2  optimizer = torch.optim.AdamW([
3      {'params': model.layer4[1].conv1.parameters(),
4       'lr': 1e-3, 'weight_decay': 0.005}, # 高敏感層
5      {'params': model.layer4[2].conv3.parameters(),
6       'lr': 8e-4, 'weight_decay': 0.007} # 次敏感層
7  ], eps=1e-5)
8
9  # 添加梯度裁剪
10 torch.nn.utils.clip_grad_norm_(
11     model.layer4[1].conv1.parameters(),
12     max_norm=2.0
13 )
```

As we can see the result, There is no sharp up and down compared with freeze layers.

However, it loses its potential to get higher accuracy, which both oscillate around that seems to overfit.



## References

1. [Code template](#)
2. [Billion-scale semi-supervised learning for image classification](#)
3. [Aggregated Residual Transformations for Deep Neural Networks](#)

4. [Deep Residual Learning for Image Recognition](#)